

Machine Intelligence 2

4.4 Locally Linear Embedding


Prof. Dr. Klaus Obermayer

Fachgebiet Neuronale Informationsverarbeitung (NI)

SS 2017

Locally Linear Embedding

Project data into the tangential (linear) space of the data manifold



- data points $\underline{x}^{(\alpha)} \in \mathbb{R}^N$
- embedded data points $\underline{u}^{(\alpha)} \in \mathbb{R}^M$, $M < N$


Locally Linear Embedding

Project data into the tangential (linear) space of the data manifold

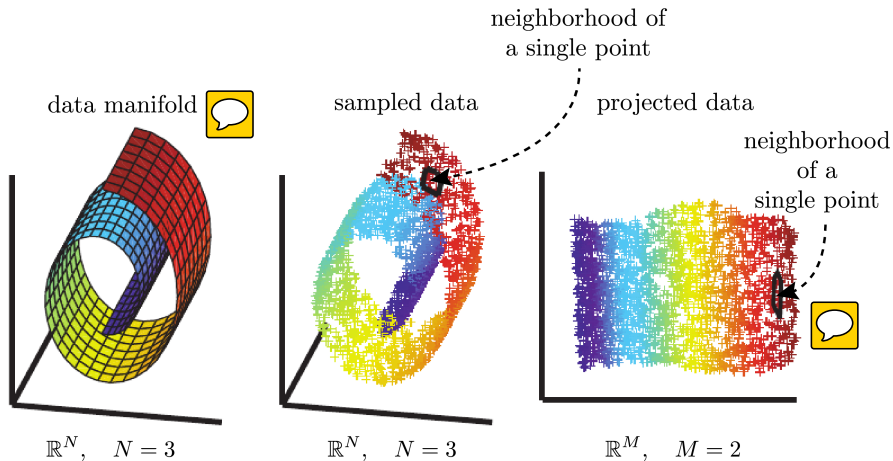
For each data point $\underline{\mathbf{x}}^{(\alpha)}$

- find the K nearest neighbors
- calculate reconstruction weights $\overline{\mathbf{W}}$
s.t. $\underline{\mathbf{x}}^{(\alpha)} \approx \sum_{\beta \in \text{KNN}(\underline{\mathbf{x}}^{(\alpha)})} \overline{\mathbf{W}}_{\alpha\beta} \cdot \underline{\mathbf{x}}^{(\beta)}$

- obtain embedding $\underline{\mathbf{u}}^{(\alpha)} \in \mathbb{R}^M$
s.t. $\underline{\mathbf{u}}^{(\alpha)} \approx \sum_{\beta \in \text{KNN}(\underline{\mathbf{x}}^{(\alpha)})} \overline{\mathbf{W}}_{\alpha\beta} \cdot \underline{\mathbf{u}}^{(\beta)}$



Locally Linear Embedding



Source: Science; Roweis, Saul 2000, modified

Step 1: find K nearest neighbors

choice: Euclidean distance

$$\beta_1^{(\alpha)} = \arg \min_{\beta} \left| \underline{\mathbf{x}}^{(\alpha)} - \underline{\mathbf{x}}^{(\beta)} \right|$$

$$\beta_2^{(\alpha)} = \arg \min_{\beta \neq \beta_1^{(\alpha)}} \left| \underline{\mathbf{x}}^{(\alpha)} - \underline{\mathbf{x}}^{(\beta)} \right|$$

$$\vdots$$

$$\beta_K^{(\alpha)} = \arg \min_{\substack{\beta \neq \beta_k^{(\alpha)}, \\ k=1, \dots, K-1}} \left| \underline{\mathbf{x}}^{(\alpha)} - \underline{\mathbf{x}}^{(\beta)} \right|$$

$$\text{KNN}(\underline{\mathbf{x}}^{(\alpha)}) = \left\{ \beta_1^{(\alpha)}, \beta_2^{(\alpha)}, \dots, \beta_K^{(\alpha)} \right\} \quad (\text{not necessarily unique})$$



linear data structure (e.g. data matrix): $\mathcal{O}(Np^2)$

k-d tree (balanced search tree): $\mathcal{O}(Np \log p)$

oneNote

Step 2: calculate reconstruction weights

minimize cost function:

$$E(\mathbf{W}) = \sum_{\alpha=1}^p \underbrace{\left\| \mathbf{x}^{(\alpha)} - \sum_{\beta=1}^p W_{\alpha\beta} \mathbf{x}^{(\beta)} \right\|^2}_{\substack{\text{reconstruct } \mathbf{x}^{(\alpha)} \text{ by its} \\ K \text{ nearest neighbors only}}} \stackrel{!}{=} \min \quad \text{s.t.} \quad \begin{aligned} W_{\alpha\beta} &= 0 \text{ if } \beta \notin \text{KNN}(\mathbf{x}^{(\alpha)}), \\ \sum_{\beta=1}^p W_{\alpha\beta} &= 1 \end{aligned}$$

weight matrix \mathbf{W} :

- sparse: (up to) K nonzero elements per row
- not symmetric: nearest neighbors of a data point can have closer neighbors

Step 2: calculate reconstruction weights

minimize cost function:

$$E(\underline{\mathbf{W}}) = \sum_{\alpha=1}^p \underbrace{\left| \underline{\mathbf{x}}^{(\alpha)} - \sum_{\beta=1}^p W_{\alpha\beta} \underline{\mathbf{x}}^{(\beta)} \right|^2}_{\substack{\text{reconstruct } \underline{\mathbf{x}}^{(\alpha)} \text{ by its} \\ K \text{ nearest neighbors only}}} \stackrel{!}{=} \min \quad \text{s.t.} \quad \begin{aligned} W_{\alpha\beta} &= 0 \text{ if } \beta \notin \text{KNN}(\underline{\mathbf{x}}^{(\alpha)}), \\ \sum_{\beta=1}^p W_{\alpha\beta} &= 1 \end{aligned}$$

optimal weights are invariant to:




- rotation $\underline{\mathbf{R}}$: $E \left[\underline{\mathbf{R}} \cdot \underline{\mathbf{x}}^{(1)}, \dots, \underline{\mathbf{R}} \cdot \underline{\mathbf{x}}^{(p)} \right] \stackrel{\text{orthog.}}{=} \underline{\mathbf{R}} E \left[\underline{\mathbf{x}}^{(1)}, \dots, \underline{\mathbf{x}}^{(p)} \right]$
- scaling $\gamma > 0$: $E \left[\gamma \underline{\mathbf{x}}^{(1)}, \dots, \gamma \underline{\mathbf{x}}^{(p)} \right] = \gamma^2 E \left[\underline{\mathbf{x}}^{(1)}, \dots, \underline{\mathbf{x}}^{(p)} \right]$
- translation $\Delta \underline{\mathbf{x}}$: $E \left[\underline{\mathbf{x}}^{(1)} + \Delta \underline{\mathbf{x}}, \dots, \underline{\mathbf{x}}^{(p)} + \Delta \underline{\mathbf{x}} \right] \stackrel{\sum_{\beta} W_{\alpha\beta}=1}{=} E \left[\underline{\mathbf{x}}^{(1)}, \dots, \underline{\mathbf{x}}^{(p)} \right]$

Step 2: calculate reconstruction weights

oneNote

minimize cost function:

$$E(\underline{\mathbf{W}}) = \sum_{\alpha=1}^p \underbrace{\left| \underline{\mathbf{x}}^{(\alpha)} - \sum_{\beta=1}^p W_{\alpha\beta} \underline{\mathbf{x}}^{(\beta)} \right|^2}_{\text{reconstruct } \underline{\mathbf{x}}^{(\alpha)} \text{ by its } K \text{ nearest neighbors only}} \stackrel{!}{=} \min \quad \text{s.t.} \quad \begin{aligned} &W_{\alpha\beta} = 0 \text{ if } \beta \notin \text{KNN}(\underline{\mathbf{x}}^{(\alpha)}), \\ &\sum_{\beta=1}^p W_{\alpha\beta} = 1 \end{aligned}$$


for each data point $\underline{\mathbf{x}}^{(\alpha)}$:

- local "covariance" matrix (symmetric & positive semidefinite) $\underline{\mathbf{C}}^{(\alpha)} \in \mathbb{R}^{K,K}$:

$$C_{jk} = (\underline{\mathbf{x}}^{(\alpha)} - \underline{\mathbf{x}}^{(\beta_j^{(\alpha)})})^T (\underline{\mathbf{x}}^{(\alpha)} - \underline{\mathbf{x}}^{(\beta_k^{(\alpha)})})$$

- solve linear system $\underline{\mathbf{C}}^{(\alpha)} \underline{\tilde{\mathbf{w}}}^{(\alpha)} = (1, \dots, 1)^T$

- rescale weights: $W_{\alpha\beta_j^{(\alpha)}} = \tilde{w}_j^{(\alpha)} / \sum_{k=1}^K \tilde{w}_k^{(\alpha)}$ to fulfill constraint

$\Rightarrow \underline{\mathbf{W}}$ contains the optimal weights with $W_{\alpha\beta} = 0$ for $\beta \notin \text{KNN}(\underline{\mathbf{x}}^{(\alpha)})$

$\Rightarrow p$ dense K -dim. linear systems have to be solved: $\mathcal{O}(pK^3)$



Step 3: obtain embedding coordinates

For any M -dimensional manifold there exist linear mappings of each local "patch" onto M -dimensional coordinates in a linear space

- linear mapping: rotation, scaling, translation
- weights $\underline{\mathbf{W}}$ can be used to optimally reconstruct the data points in the lower-dimensional embedding space

idea:

- cut N -d manifold into small patches
- "glue" them together in M -d using only rotation, scaling, translation for each patch

Step 3: obtain embedding coordinates

given $M \ll N$ and $\underline{\mathbf{W}}$: find optimal coordinates $\underbrace{\underline{\mathbf{u}}^{(1)}, \dots, \underline{\mathbf{u}}^{(p)}}_{=\underline{\mathbf{U}} \in \mathbb{R}^{M,p}} \in \mathbb{R}^M$

cost function:

$$F(\underline{\mathbf{U}}) = \sum_{\alpha=1}^p \left| \underline{\mathbf{u}}^{(\alpha)} - \sum_{\beta=1}^p W_{\alpha\beta} \underline{\mathbf{u}}^{(\beta)} \right|^2$$

equivalent quadratic form:

$$F(\underline{\mathbf{U}}) = \sum_{\alpha,\beta=1}^p g_{\alpha\beta} (\underline{\mathbf{u}}^{(\alpha)})^T \underline{\mathbf{u}}^{(\beta)}$$

where $g_{\alpha\beta} =$

see blackboard

$$= \delta_{\alpha\beta} - W_{\alpha\beta} - W_{\beta\alpha} + \sum_{\gamma=1}^p W_{\gamma\alpha} W_{\gamma\beta}$$

$\underline{\mathbf{G}} = \{g_{\alpha\beta}\} \in \mathbb{R}^{p,p}$ is symmetric and positive semidefinite

Step 3: obtain embedding coordinates

minimize cost function:



$$F(\underline{\mathbf{U}}) = \sum_{\alpha, \beta=1}^p g_{\alpha\beta} (\underline{\mathbf{u}}^{(\alpha)})^T \underline{\mathbf{u}}^{(\beta)}$$

$$\text{s.t. } \sum_{\alpha=1}^p \underline{\mathbf{u}}^{(\alpha)} = \mathbf{0}, \quad (\text{remove translation freedom})$$

$$\frac{1}{p} \sum_{\alpha=1}^p \underline{\mathbf{u}}^{(\alpha)} (\underline{\mathbf{u}}^{(\alpha)})^T = \underline{\mathbf{I}} \quad (\text{prevent trivial solutions e.g., } \underline{\mathbf{u}}^{(\alpha)} = \underline{\mathbf{0}})$$

\leadsto w.l.o.g. as $F(\underline{\mathbf{U}})$ invariant to rotation, scaling, translation

\leadsto implies that reconstruction errors for different coordinates are measured on the same scale

Step 3: obtain embedding coordinates

solution:

compute the $M + 1$ eigenvectors of $\underline{\mathbf{G}}$ with the lowest eigenvalues but discard the eigenvector $\underline{\mathbf{e}}_p = \frac{1}{p}(1, \dots, 1)^T$ with eigenvalue 0 (corresponding to translation)

$$\underline{\mathbf{U}} = \begin{pmatrix} \underline{\mathbf{e}}_{p-M}^T \\ \vdots \\ \underline{\mathbf{e}}_{p-1}^T \end{pmatrix} = \left(\underline{\mathbf{u}}^{(1)}, \dots, \underline{\mathbf{u}}^{(p)} \right) \quad (\text{solution satisfies both constraints})$$

implementation:

- store $\underline{\mathbf{W}}$ in sparse matrix format (at most $K \cdot p$ non-zero elements) and calculate $\underline{\mathbf{G}} = \left(\underline{\mathbf{I}} - \underline{\mathbf{W}}^T \right) \left(\underline{\mathbf{I}} - \underline{\mathbf{W}} \right) \in \mathbb{R}^{p,p}$
- use sparse eigenvalue solvers (eigsh)

$$\underline{\mathbf{v}} \mapsto \underline{\mathbf{G}} \cdot \underline{\mathbf{v}} = \left(\underline{\mathbf{I}} - \underline{\mathbf{W}}^T \right) \left[\left(\underline{\mathbf{I}} - \underline{\mathbf{W}} \right) \underline{\mathbf{v}} \right]$$

Summary of the LLE algorithm

parameters: K, M

- ① find K nearest neighbors $\text{KNN}(\underline{\mathbf{x}}^{(\alpha)}) = \{\beta_1^{(\alpha)}, \dots, \beta_K^{(\alpha)}\} \quad \forall \alpha = 1, \dots, p$
- ② calculate (locally invariant) reconstruction weights $\underline{\mathbf{W}}$:

$$\underline{\mathbf{C}}^{(\alpha)} \underline{\tilde{\mathbf{w}}}^{(\alpha)} = (1, \dots, 1)^T, \quad \forall \alpha = 1, \dots, p$$

$$W_{\alpha\beta_j^{(\alpha)}} = \frac{\tilde{w}_j^{(\alpha)}}{\sum_{k=1}^K \tilde{w}_k^{(\alpha)}}$$

- ③ calculate the embedding coordinates $\underline{\mathbf{U}}$:


compute the $M + 1$ eigenvectors $(\underline{\mathbf{e}}_p, \dots, \underline{\mathbf{e}}_{p-M})$ of $\underline{\mathbf{G}}$
with the smallest eigenvalues

$$g_{\alpha\beta} = \delta_{\alpha\beta} - W_{\alpha\beta} - W_{\beta\alpha} + \sum_{\gamma=1}^p W_{\gamma\alpha} W_{\gamma\beta}$$

$$\underline{\mathbf{G}} \cdot \underline{\mathbf{e}}_j = \lambda_j \underline{\mathbf{e}}_j \qquad \underline{\mathbf{U}} = \begin{pmatrix} \underline{\mathbf{e}}_{p-M}^T \\ \vdots \\ \underline{\mathbf{e}}_{p-1}^T \end{pmatrix} = (\underline{\mathbf{u}}^{(1)}, \dots, \underline{\mathbf{u}}^{(p)})$$


Example 1

- Images of a single face translated across a two-dimensional background of noise
- PCA fails to preserve the neighborhood structure of nearby images
- LLE maps the images with corner faces to the corners of its two dimensional embedding



Source: An Introduction to Locally Linear Embedding; Saul, Roweis 2001

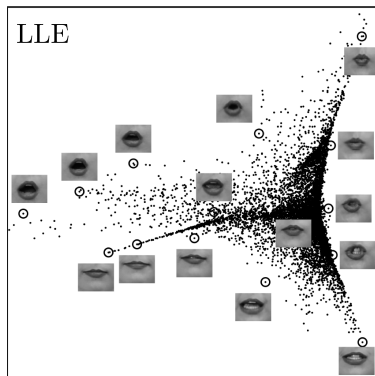
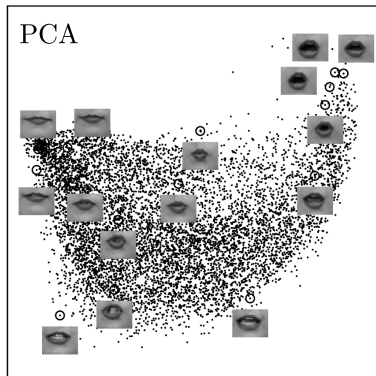
Example 2



Source: *Science*; Roweis, Saul 2000

- Images of faces mapped into an embedding space with $M = 2$.
- Bottom: Images corresponding to points along the solid line shown on the top-right.

Example 3



Source: An Introduction to Locally Linear Embedding; Saul, Roweis 2001

- Images of lips mapped into an embedding space with $M = 2$.
- Differences between the two embeddings indicate the presence of nonlinear structure in the data.

Remarks

- efficient & robust algorithm
- parameters: number K of neighbors, embedding dimension M
- convex optimization problem, standard (sparse) linear algebra methods suffice
- for $K > N$ regularization is required (singular covariance matrix $\underline{\mathbf{C}}^{(\alpha)}$)

$$\underline{\mathbf{C}}^{(\alpha)} \leftarrow \underline{\mathbf{C}}^{(\alpha)} + \varepsilon \mathbf{I}$$

- extension for pairwise data available via non-Euclidean distances $d_{\alpha\alpha'}$ in $\underline{\mathbf{C}}^{(\alpha)}$
- alternative methods available (e.g. Laplacian eigenmaps, t-stochastic neighbor embedding, isomap, multi-dimensional scaling, Kernel PCA)